

Predction Using ML

```
In [1]: # importing necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

In [2]: # Load the dataset
path = "https://raw.githubusercontent.com/AdiPersonalWorks/Random/master/student_scores%20-%20student_scores.csv"

# Read the csv file using read_csv
df = pd.read_csv(path)
```

```
In [3]: df.head()
```

Out[3]:

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

Descriptive Statistics

```
In [4]: df.describe()
```

Out[4]:

	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000


```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   Hours   25 non-null      float64
 1   Scores  25 non-null      int64
dtypes: float64(1), int64(1)
memory usage: 528.0 bytes
```

Plot the relationship between Hours and scores

```
In [6]: sns.scatterplot(x='Hours', y='Scores', data=df)
plt.xlabel('Hours studied')
plt.ylabel('Scores')
```

Out[6]: Text(0, 0.5, 'Scores')



Preparing the data

```
In [7]: # independent variable
X = df.iloc[:, :-1].values

# dependent(target) variable
y = df.iloc[:, 1].values
```

Training the model

```
In [8]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=0)
```


```
In [9]: from sklearn.linear_model import LinearRegression
lm = LinearRegression()

# Fitting the data
lm.fit(X_train, y_train)
```

Out[9]: LinearRegression()

Plotting the regression line

```
In [12]: Regression_Line = lm.coef_*X+lm.intercept_
plt.scatter(X,y)
plt.plot(X,Regression_Line);
plt.show()
```



Making Predictions

```
In [13]: # Predicting the data
predictions = lm.predict(X_test)
```

```
In [14]: # Comparing Actual vs Predicted
compare_scores = pd.DataFrame({'Actual Marks': y_test, 'Predicted Marks': predictions})
compare_scores
```

Out[14]:

	Actual Marks	Predicted Marks
0	20	16.884145
1	27	33.732261
2	69	75.357018
3	30	26.794801
4	62	60.491033

Evaluating the model

```
In [15]: from sklearn.metrics import mean_absolute_error

mae = mean_absolute_error(y_test, predictions)
print('Mean Absolute Error: ', mae)
```

Mean Absolute Error: 4.183859899002982

Testing the model by giving my values

```
In [16]: hours = [9.25]
predicted_score = lm.predict([hours])
print("Predicted Score = {}".format(round(predicted_score[0],3)))
```

Predicted Score = 93.692

In []: